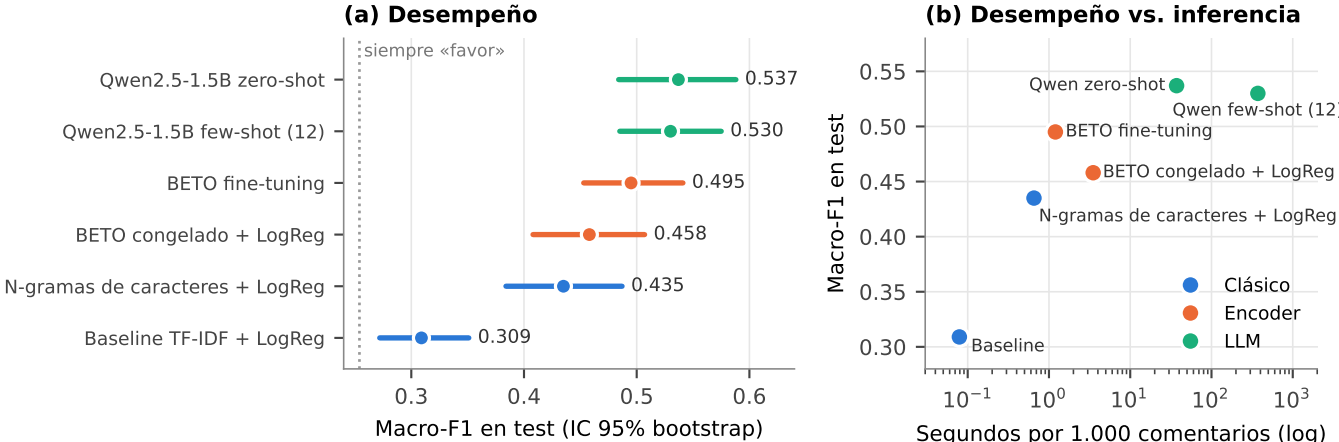
(a) Desempeño
siempre «favor»
Qwen2.5-1.5B zero-shot
Qwen2.5-1.5B few-shot (12)
BETO fine-tuning
BETO congelado + LogReg
N-gramas de caracteres + LogReg
Baseline TF-IDF + LogReg
0.537
0.530
0.495
0.458
0.435
0.309
0.3
0.4
0.5
0.6
Macro-F1 en test (IC 95% bootstrap)
(b) Desempeño vs. inferencia
Macro-F1 en test
0.55
0.50
0.45
0.40
0.35
0.30
Qwen zero-shot
Qwen few-shot (12)
BETO fine-tuning
BETO congelado + LogReg
N-gramas de caracteres + LogReg
Baseline
Clásico
Encoder
LLM
10^{-1}
10^{0}
10^{1}
10^{2}
10^{3}
Segundos por 1.000 comentarios (log)